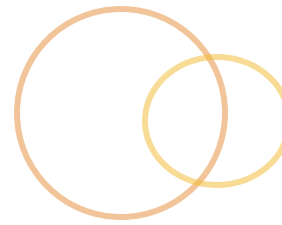


Java NIO

This is not your Father's Matrix...



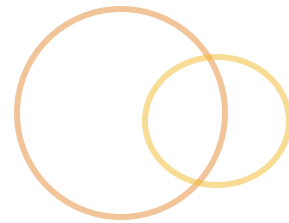
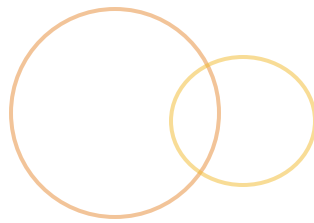
Presentation Topics



In this presentation, we will cover:

- Introduction to Java I/O
- Introduction to Java NIO
- Advanced NIO


Objectives




When we are done, you should be able to:

- Describe the differences between Java I/O and Java NIO
- Identify the key components of NIO
- Write a basic file reader using NIO
- Write a basic network handler using NIO

Introduction to I/O



A Quick Review

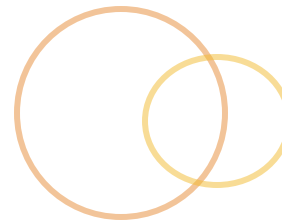
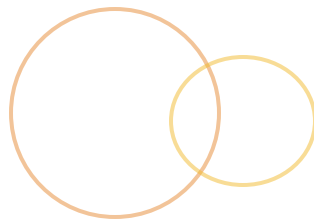




Intro to I/O

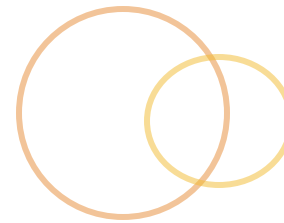
- Stands for input / output
- Input / output interface between application and operating system
- Typically seen as stream of data
- Streams filled and emptied with bytes

Java I/O



- Been around since JDK 1.0
- Found in `java.io`
- Two fundamental stream types
 - Binary
 - `InputStream`
 - `OutputStream`
 - Single byte focus
 - Character
 - `Reader`
 - `Writer`
 - Single character focus

Java I/O [cont.]



- ◎ Hides I/O details
 - ◎ Implemented in layered approach
 - ◎ Abstracts OS
- ◎ Supports stream chaining
 - ◎ Form of Decorator pattern
 - ◎ Convert stream into “higher-level” I/O construct
- ◎ Integrated with networking capabilities

Java I/O Example



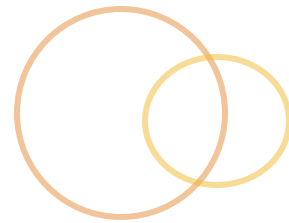
```
1 package examples.io;
2
3 import java.io.*;
4
5 public class CopyFileIO {
6
7     public static void main(String[] args) {
8         File original = new File("/tmp/pic.jpg");
9         File copy = new File("/tmp/pic_copy.jpg");
10        int fileLength = (int) original.length();
11        InputStream originalStream = null;
12        OutputStream copyStream = null;
13
14        try {
15            originalStream = new FileInputStream(original);
16            copyStream = new FileOutputStream(copy);
17            byte [] contents = new byte[fileLength];
18            originalStream.read(contents);
19            copyStream.write(contents);
20        } catch (IOException ioe) {
21            ioe.printStackTrace();
22        } finally {
23            try{
24                originalStream.close();
25            } catch (IOException ioe) {}
26            try{
27                copyStream.close();
28            } catch (IOException ioe) {}
29        }
30    }
31 }
```


Introduction to Java NIO

Channel the Stream into a Buffer




Overview of NIO



- Introduced with JDK 1.4
- Shift to block-oriented I/O
- Supports non-blocking I/O facilities
- `java.io` re-implemented using NIO

Motivations for NIO



- ◎ Provide high-speed, high-throughput I/O
 - ◎ Avoid having to use native code
 - ◎ Available through Java platform
- ◎ Support asynchronous interactions

I/O and NIO Differences



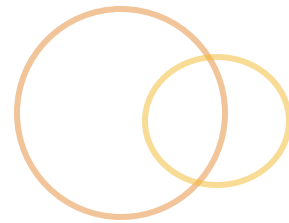
○ I/O

- Stream based; one byte at a time
- Blocking
- Easy to “build” up with chains
- Slow

○ NIO

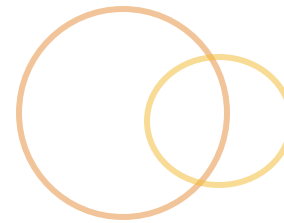
- Block based; produces / consumes block of data in one operation
- Non-blocking
- Not as elegant
- Fast

NIO Components



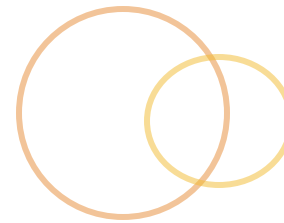
- Central Themes of NIO
 - Buffers
 - Channels
 - Selectors and selection keys
 - Charsets
- Found in 3 primary packages
 - `java.nio`
 - `java.nio.channels`
 - `java.nio.charset`

NIO Buffers



- Data container
 - Conceptually an array of byte
 - Provides structured access to data
 - Represents block
- Used by Channels for read/write operations
 - Reads fill a Buffer
 - Writes drain a Buffer
- Tracks read/write interactions

java.nio.Buffer



- Abstract class

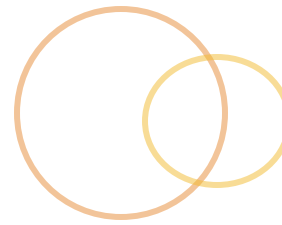
- Parent for all other `Buffers`
- Supports only primitive data elements

- Basic buffer characteristics:

- **Position** - index representing where should read / write
- **Limit** - value representing first element that should not be read / written
- **Capacity** - value representing number of elements buffer contains
- **Mark** – positional memory

$0 \leq \text{mark} \leq \text{position} \leq \text{limit} \leq \text{capacity}$

Buffer Management



- ◉ Marking - `mark ()` : `Buffer`
 - ◉ Sets a mark in the buffers at current position
- ◉ Resetting - `reset ()` : `Buffer`
 - ◉ Returns position to mark
- ◉ Clearing - `clear ()` : `Buffer`
 - ◉ Clears buffer
 - ◉ Position set to 0; limit set to capacity; marks removed
 - ◉ Makes buffer ready for reads; Call before filling buffer
 - ◉ Does not “clear” the data

Buffer Management [cont.]



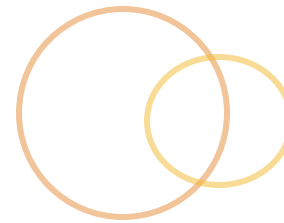
- ◎ Flipping - `flip()` : `Buffer`
 - ◎ Flips buffer
 - ◎ Limit set to current position; position set to 0; marks removed
 - ◎ Makes buffer ready for writes; Call before emptying buffer
- ◎ Rewinding - `rewind()` : `Buffer`
 - ◎ Rewinds buffer
 - ◎ Sets position to 0; removes marks
 - ◎ Makes buffer ready for re-reading info

Buffer Implementations



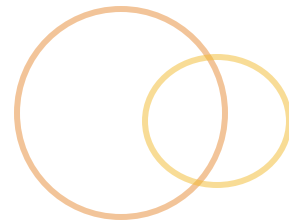
- ◎ Buffer implementation for every primitive . . .
 - ◎ `ByteBuffer`, `ShortBuffer`, `IntBuffer`, `LongBuffer`
 - ◎ `FloatBuffer`, `DoubleBuffer`
 - ◎ `CharBuffer`
 - ◎ . . . except `boolean`
- ◎ Creating Buffers
 - ◎ Allocate memory
 - ◎ Wrap existing collection

NIO Channels




- Connection to something that can do I/O
 - Files, Sockets, etc
 - Similar to streams in `java.io`
 - All NIO goes through channels
- Don't directly read / write to `Channel`
 - All data moved through `Channel` using `Buffer`
 - Bi-directional I/O
- Cannot create an instance; derive from I/O entity

java.nio.Channel




- Super-interface for all NIO channels
 - Almost tag-like interface
 - Provides two basic operations: `close` and `isOpen`
- Channel realization found in **large** type hierarchy
 - `FileChannel`
 - `ByteChannel`, `InterruptibleChannel`
 - `ScatteringChannel`, `GatheringChannel`
 - etc.
- Reading / Writing not available on all channels
 - `ReadableByteChannel`
 - `WritableByteChannel`

FileChannel



- ◎ `java.nio.channels.FileChannel`
 - ◎ Channel used to access (r/w) to files
 - ◎ Similar to `RandomAccessFile` in terms of functionality
- ◎ Must be derived from:
 - ◎ `FileInputStream`
 - ◎ `FileChannel fc = fis.getChannel()`
 - ◎ Only represents a readable channel
 - ◎ `FileOutputStream`
 - ◎ `FileChannel fc = fos.getChannel()`
 - ◎ Only represents a writable channel
 - ◎ Or, utility methods in `Channels` class

FileChannel [cont.]



Full featured file support

- ◎ File locking
 - ◎ lock
 - ◎ tryLock
 - ◎ FileLock class
- ◎ Optimized data transfer
 - ◎ transferFrom
 - ◎ transferTo
- ◎ Truncating - truncate

Basic NIO File I/O Example



- ◎ Copy contents of one file into another
- ◎ Steps involved
 - ◎ Create data holder for Channel
 - ◎ Create Channel for old File and new File
 - ◎ Read data from old File into Buffer
 - ◎ Writing data from Buffer into new File

NIO File Copy Example



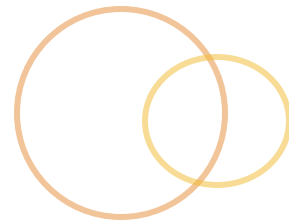
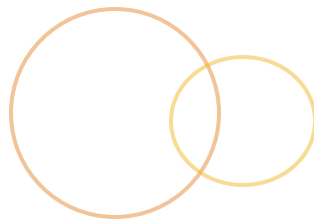
```
1  package examples.nio;
2
3  +import ...
9
10 public class CopyFileNIO {
11
12     public static void main(String[] args) throws IOException {
13         File readFile = new File("/tmp/pic.jpg");
14         File writeFile = new File("/tmp/pic_copy.jpg");
15
16         ByteBuffer fileBuffer = ByteBuffer.allocate(256);
17
18         FileInputStream fis = new FileInputStream(readFile);
19         FileChannel origChannel = fis.getChannel();
20         FileOutputStream fos = new FileOutputStream(writeFile);
21         FileChannel copyChannel = fos.getChannel();
22
23         int bytesRead = origChannel.read(fileBuffer);
24         while(bytesRead != -1) {
25             fileBuffer.flip();
26             copyChannel.write(fileBuffer);
27             fileBuffer.clear();
28             bytesRead = origChannel.read(fileBuffer);
29         }
30
31         copyChannel.close();
32         origChannel.close();
33         fos.close();
34         fis.close();
35     }
36 }
```


NIO LAB: Write a File Copier



- Implement a file copier using a `FileChannel`, and a `ByteBuffer`. Refer to the example for guidance.

ByteBuffer



- Support features you expect of buffer
 - Allocating with specific size
 - Wrapping array as buffer
 - Slicing a buffer into two
 - Making it read-only

ByteBuffer Configurations



- ⦿ Supports three types of configurations
 - ⦿ Indirect - Fast
 - ⦿ Read / writes stored in intermediary buffer before I/O operations
 - ⦿ Default configuration
 - ⦿ Direct - Faster
 - ⦿ Performs native I/O operations directly on buffer
 - ⦿ Created using `allocatedDirect`
 - ⦿ May exist outside of GC heap

ByteBuffer Configurations [cont.]






- Supports three types of configurations [cont.]
 - Memory Mapped - Fastest*
 - Map portion of file to physical memory
 - Created using `FileChannel.map`
 - Returns specific type of `ByteBuffer`
 - `MappedByteBuffer`
 - Mapping based on `FileChannel.MapMode`
 - `READ_ONLY`
 - `READ_WRITE`
 - `PRIVATE`
 - Changes dependent on OS
 - *Fastest when dealing with large files

LAB : Write a Mapped Byte Buffer

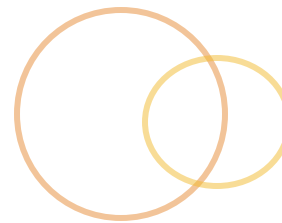


- Write a file copier that uses a direct byte buffer to perform copying.
- Do a time based test against the direct and indirect file copy mechanism.
- Write one more file copy using standard I/O.
- Re-run the time test again.
- Which one wins? By how much?

Advanced NIO



Advanced NIO Topics



- Scattering / Gathering
- Asynchronous I/O

Scattering and Gathering



Scattering

- Reading bytes from channel into multiple buffers
- Requires `ScatteringByteChannel`

Gathering

- Writing bytes to channel from multiple buffers
- Requires `GatheringByteChannel`

Good when you are reading / writing mixed, fixed-length content

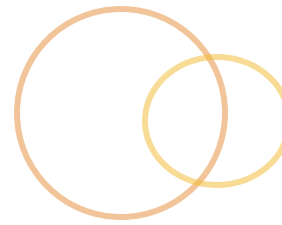
- HTTP Request contains 3 sections
- Could have three buffers; one for each section

Asynchronous I/O

No more lines . . .

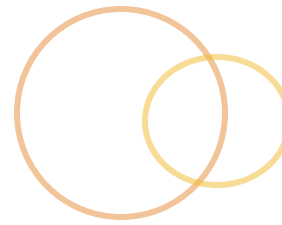


Asynchronous I/O



- Reading / writing data without blocking
 - No waiting for data on read or for access on write
 - Can use single thread to do **ALL** I/O operations
- Operations occur as result of notification
 - Event-like, notification-based system
 - Register “interest” in specific events
 - Associate “registrations” with Channels
- Relies on:
 - `SelectableChannels` - configured to support NIO
 - `Selector` - event medium
 - `SelectionKey` - registration identification
 - `Handler` - entity that processes “events”

SelectableChannel



- ◉ Channel that supports selection
- ◉ Safe for multi-threaded interactions
- ◉ Supports two modes:
 - ◉ Blocking
 - ◉ Every I/O operations blocks until completes
 - ◉ Default configuration
 - ◉ Non-blocking
 - ◉ No I/O blocks
 - ◉ Configured through `configureBlocking` method
 - ◉ Must be configured to non-blocking before registration with `Selector`

SelectableChannel [cont.]



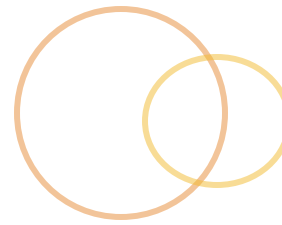
- Selection functionality provided through registration
 - Registration relies on:
 - Selector - event facilitator
 - I/O op codes - event “ids”
 - Performed through:
 - `public SelectionKey register(Selector sel, int ops)`
 - `public SelectionKey register(Selector sel, int ops, Object att)`

SelectableChannel [cont.]



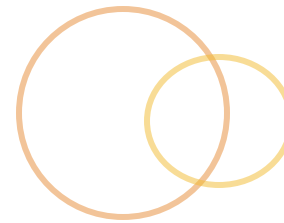
- Selection functionality provided through registration
 - Registration represented as `SelectionKey`
 - Valid until key is canceled
 - Or when channel closes
 - Or when `Selector` closes
- Common implementations:
 - `ServerSocketChannel`
 - `SocketChannel`

ServerSocketChannel



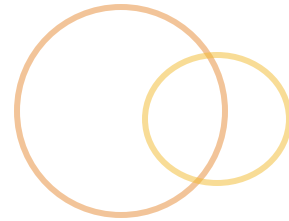
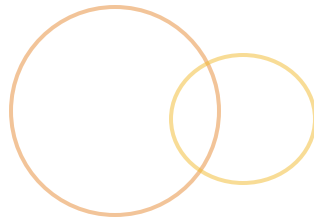
- Channel representation of a `ServerSocket`
 - All `ServerSocketChannel` objects have a associated `ServerSocket`
 - Not all `ServerSocket` have an associated `ServerSocketChannel`
- Creating a `ServerSocketChannel`

```
Selector selector = Selector.open();
ServerSocketChannel channel = ServerSocketChannel.open();
ServerSocket socket = channel.socket();
socket.bind(XXXX);
//configure non-blocking
channel.configureBlocking(false);
channel.register(selector, SelectionKey.OP_ACCEPT);
```



- ◉ Channel representation of a Socket
 - ◉ All `SocketChannel` objects have a associated `Socket`
 - ◉ Not all `Socket` have an associated `SocketChannel`
- ◉ Creating a `SocketChannel`

```
Selector selector = Selector.open();  
SocketChannel channel = SocketChannel.open();  
Channel.connect(channel, xxxx);  
Socket socket = channel.socket();  
//configure non-blocking  
channel.configureBlocking(false);  
channel.register(selector, SelectionKey.OP_ACCEPT);
```



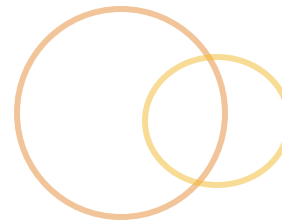
- ◎ Functions as an event medium facilitator
 - ◎ Not true event listener / event handler
 - ◎ More of a poll-based versus notification-based mechanism
 - ◎ Binds channel, I/O events, and handler
- ◎ Basic operation of Selector
 - ◎ Receives registration interest in I/O related events
 - ◎ Channel notifies Selector when “event” occurs
 - ◎ Poll Selector for “events”

Selector [cont.]



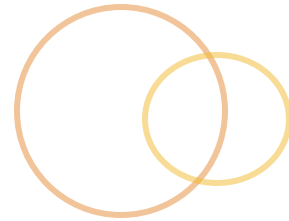
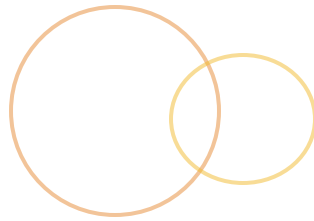
- ◎ `java.nio.channels.Selector`
 - ◎ Created using `open` call; functions as Factory method
 - ◎ Associated with `SelectableChannel`
 - ◎ Contains `SelectionKey` sets
 - ◎ Considered multi-threaded safe
- ◎ Two main functions:
 - ◎ Registration - represented by `SelectionKey`
 - ◎ Selection - getting subset of `SelectionKey` representing channels ready for I/O operations

SelectionKey



- Represents registration of Selector with Channel
 - New key is created for each registration
 - Key stays valid until canceled
- Used as “notification” mechanism when event occurs
 - “Event id” represented as op code
 - 4 predefined op-codes:
 - OP_ACCEPT
 - OP_CONNECT
 - OP_READ
 - OP_WRITE
- Can contain attachment accessible by “handler”

Handler



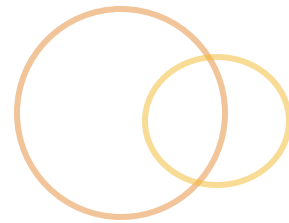
- ⦿ No “handler” interface
- ⦿ Three main functions:
 - ⦿ Select ready channels
 - ⦿ `select()` : `int` - synchronous
 - ⦿ `select(long timeout)` : `int` - synchronous
 - ⦿ `selectNow()` : `int` - asynchronous
 - ⦿ Retrieves selected channels
 - ⦿ `selectedKeys()` : `Set<SelectionKey>` - result not thread safe
 - ⦿ Processes results
 - ⦿ `channel()` : `SelectableChannel` - “ready” channel
 - ⦿ read / write
 - ⦿ remove key from selection
- ⦿ Typically written as loop

Asynchronous I/O Example



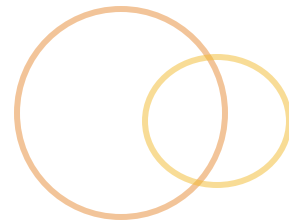
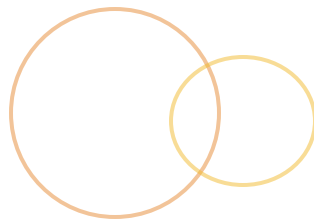
- Traditionally, needed one thread per connection to handle “many requests”
 - Can adopt NIO strategy
 - One thread for “accepts”
 - One thread for request processing
 - Three key components
 - Main - application launcher
 - ServerSocketThread - accepts connections
 - SocketThread - processes requests

NIO Charsets



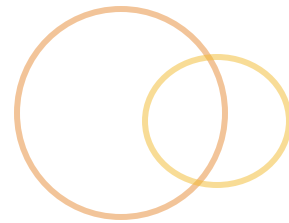
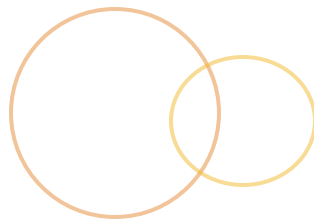
- “named mapping between sequences of 16-bit Unicode characters and sequences of 8-bit characters”
- NIO supports character set mappings on channels
 - Not typically used by developer
 - Defined in terms of:
 - `CharsetEncoder` - encodes sequence of characters
 - `CharsetDecoder` - decodes sequence of characters
- Found in `java.nio.charset`

Summary



- ◎ Java NIO expands I/O to a channel model
- ◎ Channels are like streams of data that support bi-directional I/O operations
- ◎ Channels rely on Buffers to read / write data
- ◎ SelectableChannels support asynchronous I/O

NIO Lab 2



- Description: Convert the sitemap utility that was created in the advanced threading chapter into a asynchronous NIO application. The resulting application will function as an asynchronous network client, similar to how a browser works when you surf the web.
- Duration: 1 hour